

# End-to-End Inference Latency vs. Clinical Faithfulness on 62k Corpus

